

MOBILITY PRICING INTELLIGENCE: COST PREDICTION & PRICING DRIVER ANALYSIS

Inventory Planning, Pricing Dynamics
& Operational Performance



TOOLS USED

Python, Pandas, Numpy, Matplotlib, Plotly &
ML libraries

PRESENTED BY

Tejas Jadhav





TABLE OF CONTENTS

Overview & Business Objective	3
Data Understanding & EDA	8
Data Preparation & Feature Engineering	22
Model Development & Experimentation	29
Model Evaluation & Comparison	34
Explainability & Final Model	44
Business Insights & Recommendations	49
Conclusion	54



Overview & Business Objective



1.1 PROJECT OVERVIEW

Retail businesses operate in a dynamic environment where demand is influenced by factors such as historical sales, inventory, pricing, promotions, seasonality, and region.

Accurate demand estimation helps businesses maintain optimal inventory, reduce stockouts and excess stock, and improve procurement and operational planning.

The Retail Demand Intelligence project uses historical retail data and machine learning to predict demand and identify the key factors influencing it.


The project combines:

- Data understanding and cleaning
- Exploratory Data Analysis
- Feature engineering
- Machine learning regression
- Model evaluation and comparison
- Prediction-error analysis
- Model explainability
- Business-oriented insights

The objective is not only to predict demand accurately, but also to understand why the model makes its predictions & how these insights can support better decision-making.



1.2 PROBLEM STATEMENT



Retail demand is influenced by multiple interacting variables, making manual forecasting or reliance on historical averages insufficient for many situations.

The key problem addressed in this project is:

How can historical retail data be used to build a machine learning model that accurately predicts product demand while also identifying the major factors driving demand?

The project therefore treats Demand as the target variable and uses historical, operational, pricing, promotional, environmental, and temporal information as predictive features.



1.3 PROJECT OBJECTIVES

The major objectives of the project are:

- Develop a machine learning-based retail demand prediction system.
- Understand the underlying structure and characteristics of the dataset.
- Identify important patterns through exploratory data analysis.
- Create meaningful temporal, inventory, pricing, and historical features.
- Train and compare multiple regression models.
- Evaluate model performance using appropriate regression metrics.
- Investigate the largest prediction errors.
- Identify the most influential demand-driving features.
- Identify limitations and possible improvements for future development.



1.4 PROJECT SCOPE

The project follows a structured workflow covering data analysis, demand modeling, model explainability, and business interpretation to understand demand patterns, evaluate predictive performance, identify key demand drivers, and derive actionable retail insights.

01

Data Analysis

- Dataset inspection
- Data quality assessment
- Distribution analysis
- Relationship analysis

02

Demand Modeling

- Feature preparation
- Regression modeling
- Model comparison
- Model evaluation

03


Explainability

- Feature importance analysis
- Original-feature aggregation
- Investigation of large prediction errors

04

Business Interpretation

- Business Interpretation
- Inventory implications
- Pricing and promotion observations
- Seasonal and weather-related effects
- Recommendations for retail planning



Data Understanding & Exploratory Data Analysis



2.1 DATA OVERVIEW

The dataset provides a comprehensive view of retail operations across stores and products, covering inventory levels, sales activity, ordering patterns, pricing, discounts, promotions, weather conditions, seasonality, and demand. These variables provide the foundation for analyzing demand behavior, identifying operational patterns, and developing predictive models for retail demand.

Data Source: Kaggle

Dataset Size

76,000

Records

16

Original Features

Purchase Diversity

4

Regions

5

Stores

20

Products



DATA OVERVIEW

Below is a detailed description of the feature set:

Dataset Features	Type	Feature Description
Date	Date / Time	Date of the observation
Store ID	Categorical	Unique identifier of the store
Product ID	Categorical	Unique identifier of the product
Category	Categorical	Product category
Region	Categorical	Geographic region of the store
Inventory Level	Numerical (Discrete)	Inventory available at the time of observation
Units Sold	Numerical (Discrete)	Number of units sold
Units Ordered	Numerical (Discrete)	Number of units ordered
Price	Numerical (Continuous)	Selling price of the product
Discount	Numerical (Discrete)	Discount applied to the product
Weather Condition	Categorical	Weather condition during the observation
Promotion	Binary	Indicates whether a promotion was active
Competitor Pricing	Numerical (Continuous)	Price of the comparable product offered by competitors
Seasonality	Categorical	Seasonal classification of the observation
Epidemic	Binary	Indicates whether an epidemic-related condition was present
Demand	Numerical (Discrete)	Target variable representing product demand

2.2 TARGET VARIABLE

The target variable selected for prediction is:

Demand

Demand represents the quantity that the model attempts to estimate based on the available information

The target has:

104.32

Mean

100

Median

4

Minimum

430

Maximum

46.96

Standard Deviation

The distribution is concentrated around approximately 70-130 units, with a noticeable right tail extending toward higher demand values.




2.3 DATA QUALITY & PREPARATION OVERVIEW

Before modeling, the dataset was examined for:

- Data types
- Missing values
- Duplicate observations
- Categorical variables
- Numerical variables
- Temporal information
- Potential outliers
- Feature relationships

The *Date* variable was converted into usable temporal components, while categorical variables were prepared for machine learning using one-hot encoding.






2.4 Exploratory Data Analysis



Historical ride costs are broadly distributed across a wide range, with the highest concentration occurring in the mid-range.

\$2,135

Average Cost per Service



Key observations

- Historical ride costs span approximately \$25 to \$836, with the highest concentration of observations occurring roughly between \$150 and \$500.
- The distribution is reasonably broad, with observations extending across both lower-cost and higher-cost rides rather than being concentrated around a single narrow price range.


Business Insights

- Most retail activity occurs within a relatively predictable baseline demand range, which can assist with normal inventory planning.
- The long right tail indicates occasional high-demand events, highlighting the importance of maintaining sufficient safety stock and identifying demand spikes.

Groceries dominate the dataset, creating stronger representation of grocery demand

\$1,251

Typical Cost per Resource (Median)



Key observations

- Groceries has the largest representation in the dataset, with approximately 30,000 observations.
- Electronics has the smallest representation, while Furniture, Clothing, and Toys fall between these categories.


Business Insights

- Overall dataset patterns may be influenced more strongly by the highly represented Grocery category.
- Category-specific forecasting and inventory strategies can help account for differences in product behavior.

North dominates regional representation, while the remaining ashow balanced coverage

\$2,147

Average Spending per Region



Key observations


- The North region has substantially more observations than the other regions.
- South, East, and West have relatively similar representation.

Business Insights

- Regional comparisons should account for the unequal number of observations.
- Regional segmentation can help retailers identify differences in demand behavior and plan inventory accordingly.

Retail demand fluctuates significantly over time, revealing recurring periods of elevated activity

Expected Ride Duration vs Historical Ride Cost



Key observations

- Daily demand demonstrates considerable fluctuations throughout the observation period.
- Several recurring periods of lower and higher demand are visible, including elevated demand around parts of 2023.

Business Insights

- Retail demand should be treated as a time-dependent phenomenon rather than a constant value.
- Dynamic inventory and procurement planning can help businesses respond to expected periods of higher or lower demand.

Winter dominates seasonal representation, while the remaining seasons maintain relatively balanced coverage



Key observations

- Winter contains the largest number of observations.
- Spring, Summer, and Autumn have relatively similar representation.

Business Insights

- Seasonal representation should be considered when interpreting seasonal patterns.
- Incorporating seasonality into demand prediction allows the model to capture recurring changes associated with different periods of the year.

Demand is strongly linked to sales activity, while pricing factors show limited linear influence

52.03%

Typical CPU Utilization (Median)



Key observations

- Demand has a strong positive correlation with Units Sold (0.83).
- Demand has a moderate positive correlation with Units Ordered (0.51) and a weaker relationship with Inventory Level (0.13).

Business Insights


- Historical sales and ordering activity provide valuable information for understanding demand.
- Price and competitor pricing have very weak linear correlations with Demand in this dataset, suggesting that pricing effects may depend on interactions with other factors.

(Important modelling note: The strong relationship between Demand & Units Sold is useful for exploratory analysis, but contemporaneous Units Sold needs to be interpreted carefully in the forecasting scenario because future sales would not necessarily be known at prediction time)

Higher discounts do not consistently translate into higher demand, indicating a more complex pricing relationship

51.63%

Typical Memory Utilization (Median)




Key observations

- Demand varies substantially across every discount level.
- There is no obvious simple linear relationship where increasing discount consistently produces increasing demand.

Business Insights

- Larger discounts do not automatically guarantee proportionally higher demand.
- Discount strategies should therefore be evaluated alongside product characteristics, inventory, promotions, seasonality, and other demand drivers.

Sunny conditions are associated with higher average demand, while adverse weather corresponds to lower demand



Key Observations

- Sunny conditions have the highest average demand, at approximately 115 units.
- Rainy and Snowy conditions have lower average demand, around 94–95 units.

Business Insights

- Weather conditions can contribute to differences in retail demand and should therefore be considered during forecasting.
- Businesses may benefit from adjusting inventory planning when weather patterns indicate potential changes in demand.